

SRI KRISHNA KALYAN BANDARU

Data Scientist | Adelaide, SA, Australia | Full Working Rights (485 Visa)

Phone: 0430 037 548 | Email: srikrishnakalyan7@gmail.com

Portfolio: srikrishnakalyan7.github.io | GitHub: github.com/SriKrishnaKalyan7

PROFESSIONAL SUMMARY

Recent Master of Artificial Intelligence and Machine Learning graduate from the University of Adelaide, with a first-class engineering background and sixteen months supporting users on a live enterprise system. Comfortable across the full path from raw data to a result someone can act on: cleaning and transforming messy operational data, building and validating models, and presenting findings to people without a technical background. Work to date includes a clinical prediction pipeline validated across patients, a root cause investigation that quantified the revenue lost to a service failure, and first-level technical support for staff on a 24/7 roster. Keen to apply this to virtual and in-home care, where the analysis has a direct effect on patient outcomes.

SKILLS

Programming & Data: Python (pandas, NumPy, scikit-learn, PyTorch), SQL (BigQuery, MySQL), R (statistical modelling and inference), Git, data cleaning and transformation, imputation, feature engineering

Statistics & Modelling: Predictive modelling, cross-validation, statistical modelling and inference, probability and statistics, experiment and A/B test design review, model benchmarking

Visualisation & Reporting: Power BI, Tableau, Matplotlib, Seaborn, Streamlit, Microsoft Excel (pivot tables, what-if modelling), dashboard development

Systems & Support: SAP production support, first-level troubleshooting, incident triage and escalation, shift handover and knowledge-base documentation

Platforms & Tools: Azure, AWS, Docker, FastAPI, Linux, Agile/Scrum, CI/CD

PROFESSIONAL EXPERIENCE

Data & Systems Analyst, Research Placement

Feb 2025 – Aug 2025

Auto-ID Lab, University of Adelaide / Adelaide, Australia

- Investigated repeated failures in an autonomous navigation test programme and traced the cause to poor feature tracking in low-texture environments, measuring about 5 tracked points per frame against a documented minimum of 15 to 20, then validated a replacement method that raised it to roughly 2,000
- Trained and benchmarked 5 object-detection configurations on a 28,000-instance labelled dataset, quantifying the trade-off between accuracy and processing cost so the team could select a model that fitted its hardware budget
- Analysed performance data across simulated test runs to document 4 recurring failure modes, benchmarked 3 candidate algorithms against them, and wrote up recommendations the team used to set its next direction
- Migrated the core research platform, 5 packages in total, from ROS1 to ROS2 after the original stopped building on current infrastructure, and built 3 simulation environments so changes could be tested safely

Machine Learning Engineer Intern

Aug 2024 – Feb 2025

Proto Corp Pvt Ltd / Remote, India

- Built and quality-checked a 3,508-image labelled dataset across 5 defect classes using polygon segmentation, exported in two formats so a single labelled set served both model families without duplicate annotation work
- Trained and compared 4 segmentation models on a 519-image, 1,514-instance validation set and recommended YOLOv8m-seg, which scored more than twice the accuracy of Mask R-CNN (0.259 against 0.122) at around half the processing cost
- Found that the 3 leading models scored within 0.007 of each other despite a 43% spread in processing cost, showing the limit sat in the architecture rather than the tuning, and recommended a move to Vision Transformers that the team adopted
- Built a FastAPI backend and Streamlit interface so testers and developers without direct model access could run any of the 3 trained models themselves and export structured results, removing a bottleneck on the team

Project Engineer, Production Support

May 2018 – Sept 2019

Wipro Technologies / India

- Provided first-level support to users of a live SAP environment on a 24/7 rotating roster, working through several hundred processing errors a shift and diagnosing each one well enough to resolve it or escalate it to the right specialist team
- Tracked recurring and high-volume errors and escalated them with the supporting detail, so fixes addressed the cause instead of clearing the same failure again each shift
- Reported error volumes, resolutions and open escalations at the end of every shift, giving the incoming team and the client a consistent record of system health
- Took over coordination of a 5 person support team as senior colleagues moved to other projects, owning the roster and acting as the escalation point to management

PROJECTS

Clinical Predictive Analytics, Alzheimer's Disease Diagnosis

Python, Statistical Validation

- Built a predictive pipeline reaching 81% accuracy in identifying Alzheimer's progression from noisy neurological signal data
- Validated with leave-one-patient-out cross-validation and an AUC of 0.74, so the result held across patients rather than within them, and reported both figures together so the limits of the model were visible alongside its headline accuracy

Service Failure Investigation, Olist Marketplace

SQL, Power BI, Excel

- Outcome scores had been falling for two months and one group was assumed to be at fault. Built a bank of 13 SQL queries across 5 stakeholder teams and traced the cause to a delivery failure concentrated in a single region.
- Quantified the impact: late orders scored 1.7 stars lower on average, around 2,500 orders and R\$350K of revenue were affected, and 54% of those reviews were one star
- Reviewed an A/B test read-out and found the randomisation was broken before any conclusions were drawn, then delivered the corrected findings in a 3 page Power BI dashboard and an Excel what-if model

Exploratory Data Analysis, Netflix Catalogue

Python, pandas

- Ran exploratory analysis on more than 8,800 titles using pandas, Matplotlib and Seaborn to surface distribution trends, seasonal patterns and category performance
- Handled extensive missing data with a layered groupby-mode imputation approach that kept a 202,000-row unnested dataset usable, then communicated the findings in plain business language

EDUCATION

Master of Artificial Intelligence & Machine Learning

Dec 2025

University of Adelaide | **GPA: 6.1/7** (Distinction average)

Quantitative coursework: Statistical Modelling and Inference (D), Probability and Statistics (D), Mathematical Foundations of Data Science (C)

Master of Science in Mechatronics Engineering

2022

University of Rome 'Tor Vergata' | **Grade: 106/110** | Focus: Systems Analysis, Control Systems

Bachelor of Mechanical Engineering

2018

KL University, India | **First Class with Distinction**

Mathematics coursework: 4 subjects across first, second and third year, all graded **O** (Outstanding, the highest grade awarded)